

Combining Neural Density Estimation and Computational State Space Models for Human Activity Recognition

Master Thesis

presented by
Benedikt Jakob Lasotta
Matriculation Number 1610104

submitted to the
Institute for Enterprise Systems
Dr. Stefan Lüdtke
University of Mannheim

September 2022

Contents

1	Introduction	1
1.1	Related Work	2
2	Theoretical Framework	4
2.1	Distribution Estimation with Masked Autoregressive Flow	4
2.2	Introduction to Computational State Space Models	6
3	Construction of the Hybrid Model	8
4	Experimental Evaluation	9
4.1	Settings	9
4.2	Experiments	9
4.3	Results	9
5	Conclusion	10
5.1	Summary	10
5.2	Future Work	10
A	Program Code / Resources	13
B	Further Experimental Results	14

List of Algorithms

List of Figures

List of Tables

Chapter 1

Introduction

The task of recognizing human activities from sensor data is relevant for many real world applications. Its uses range from optimizing logistics in warehouses to providing healthcare and assistance systems for elderly and disabled people [7, 3]. Consequently, human activity recognition (HAR) has been an active topic of research in recent years. Historically, methods to perform sensor-based HAR, broadly fall into either of two categories, each with their respective strengths and weaknesses: Symbolic methods as proposed in [6, 12] can employ prior domain knowledge and thus may need less training data, but often lack the capacity to learn from complex, time-series sensor data. One such approach are computational state space models (CSSMs) [6], which employ probabilistic inference on top of symbolic, rule based representations of possible actions within a domain to reason about observed activities. Contrary to this, modern, data-driven methods are capable of handling complex, multi-channel sensor data and have shown promising results in the past. An approach using a convolutional neural network (CNN) is described in [9]. A further example which is based on an encoder-only Transformer architecture is proposed in [14]. In this work the authors use a Transformer architecture to aggregate sensor data from the entire time series via the attention mechanism. Their results indicate that such a system improves prediction accuracy over models based on CNNs. Next to this, many other deep learning methods for sensor-based HAR exist. A comprehensive survey over current deep learning HAR methods and challenges is given by Chen et al. in [2]. Among the drawbacks of purely data-driven methods is, that they are oblivious to the causal structure of activities in the domain and thus may require more training data which is often hard to obtain. Thus recently, attempts have been made to integrate symbolic and data driven models. In the line of this work such approaches are referred to as hybrid models. The goal of these is to combine both approaches in order to obtain state of

the art performance while being able to incorporate prior knowledge to reduce the need for training data.

To benefit from the advantages of symbolic and data-driven methods, the main goal of this thesis is to develop a hybrid system for sensor-based HAR by combining existing components in a novel way. The aim of this hybrid system is to achieve a greater sample efficiency than current state of the art systems. That is, with a fixed amount of limited training data, the hybrid system should outperform both data-driven and purely symbolic approaches. To build this hybrid model, an implementation of computational state space models (CSSM) [6] referred to as computational causal behavior models (CCBM) will be combined with an observation model that is obtained from sensor data via neural density estimation methods. Initial experiments will be conducted with a type of normalizing flows called masked autoregressive flow (MAF) [10] since it has shown promising performance on density estimation tasks, as shown in [5]. The system will be evaluated on a data set of raw sensor data from inertial measurement units (IMUs) of people performing the task of cooking a carrot soup, eating at, and cleaning up afterwards. It is available for download from the University of Rostock [?].

Additionally, an important research goal of this thesis is to examine which component contributes in what part to the final performance. To this end, an ablation study will be conducted. For this, the performance of the symbolic approach with a simple observation model, the density estimation component with a simple prediction, the full hybrid model, and a baseline of quadratic discriminant analysis (QDA) will be compared on the data set introduced above. The results of this ablation study can be found in chapter (insert ref).

1.1 Related Work

(TODO: In this section only hybrid models are introduced as related work, since a number of symbolic/deep models are discussed in the previous section. Think about whether I should move this discussion into this section/make everything one section/leave it as is)

Intellectually closest to this thesis is the approach proposed by Rueda et al. introduced in [13]. They present the idea of using deep neural architectures as the observation model required by CSSMs. In particular the authors use a CNN to learn from time-series sensor data provided by inertial measurement units (IMUs), and predict the observed action class. The authors evaluated their model on the same carrots data set that will be used for evaluation in this thesis and found that their model performance was comparable to that of state of the art deep models at the time. In the line of this work, a similar approach will be taken. With the dif-

ference being, that the observation model is obtained by a neural density estimator like MAF. Which is subsequently combined with CSSMs to perform probabilistic inference. As the same dataset is used for evaluation, it will be interesting to compare the results with the approach above. However, many other approaches to combine deep neural architectures with symbolic reasoning have been proposed in recent literature. DeepProbLog [8] is introduced as a framework which combines neural networks and probabilistic-logical modeling via the existing language ProbLog [11]. This is done by extending ProbLog with neural predicates. These special predicates essentially compute the probability of atomic expressions in a probabilistic logic via a neural network. The evaluation of DeepProbLog shows that this framework is capable of performing symbolic reasoning as well as deep learning from examples in an end-to-end fashion. The downside of this approach is that its inference algorithms and language are ill-suited for dynamic systems, such as HAR. In [7] another hybrid architecture is proposed. This model combines high-level features extracted by a deep neural network with context information about process steps to obtain state-of-the-art HAR performance. (A very recent approach that shows promising results for HAR is DUSTIN [1] which employs knowledge infusion on top of features extracted via a neural network. That is, the model uses a CNN to extract features from time series sensor data as well as the high-level context and concatenates to this representation a set of features which are obtained by a knowledge based reasoner. Just as for CSSMs, this has the added benefit that common-sense knowledge can be introduced to the model, which means that estimated action sequences are consistent with the user-context. The evaluation of DUSTIN shows that it is able to outperform other state-of-the-art neuro-symbolic approaches while achieving a high sample efficiency at the same time.

Chapter 2

Theoretical Framework

2.1 Distribution Estimation with Masked Autoregressive Flow

The task of estimating the joint probability distribution $p(x)$ from a set of examples $\{\mathbf{x}_n\}_{n=1}^N$ is a central aspect in many machine learning applications. In the line of this work it is of great interest to obtain density estimates for the sensor data as this can be used for the observation model. This approach allows greater flexibility than using a simple multivariate normal distribution for all actions of a given class, which was used in the original feasibility study on CSSMs [6]. The following paragraphs will introduce some preliminary knowledge and terminology that is required to follow the remainder of this thesis.

Let \mathbf{x} be an example with D variables x_1, x_2, \dots, x_D . The chain rule of probability states that any joint distribution of variables can be decomposed into a product of joint conditional probabilities, where variable x_d only depends on the prior variables $x_{1:d-1}$. According to this rule, the joint distribution $p(x)$ can be expressed as

$$p(x) = \prod_{d=1}^D p(x_d | x_{1:d-1}). \quad (2.1)$$

This property can be exploited in order to estimate the joint distribution. To perform this task, many methods exist. In the line of this work the focus will be on methods based feed-forward neural networks which fulfill certain properties that allow correct estimation of probability densities. In particular this section will briefly introduce masked autoregressive flow (MAF) [10] as a model of performing neural density estimation from a set of examples. MAF can be thought of

as stacking multiple layers of autoregressive models on top of each other to obtain greater modeling flexibility. In the case of this work each individual MAF layer is a masked autoencoder for distribution estimation (MADE) [4]. A MADE is an autoencoder, that is a feed forward neural network which takes as input some D -dimensional vector \mathbf{x} and learns a compressed hidden representation $h(x)$ from which a reconstruction \hat{x} of input x is obtained. It is possible to use an autoencoder for density estimation if each partial output \hat{x}_d only depends on the previous variables x_1 to x_{d-1} and computes the probability of x_d given these previous variables. Note, that these conditionals are the factors of the joint distribution in formula 2.1 and thus their product gives the joint distribution $p(x)$. If each output unit only depends on the previous variables the model fulfills the so called *autoregressive property*. To ensure that a MADE has this property, binary masking matrices are multiplied element wise to the encoder and decoder weight matrices of the autoencoder. These masking matrices have the value 1 where connections should be retained and 0 where they should be dropped. This guarantees that variables x_1 to x_{d-1} are connected to the output \hat{x}_d , while the remaining variables x_d to x_D are dropped this process is explained in further detail in [4].

A main disadvantage of MADE is that the order of the conditionals has a large impact on the quality of the learned density. With the correct order it may be possible to learn the density perfectly while for other orders this is not the case. To remedy this, MAF models the random numbers which are used by the MADE in layer k with the model in layer $k + 1$ and so on, until the random numbers of the model in the last layer are modeled by a standard Gaussian. Although each individual MADE unit has unimodal conditionals the authors of [10] argue that such a MAF can learn multimodal conditionals. According to the authors, this added flexibility improves model fit. The central idea behind MAF is, that the distribution to be learned is viewed as a transformation of a base density $\pi_u(u)$ to the target distribution via a differentiable and invertible transformation f . The base density should be simple in the sense that $\pi_u(u)$ can be easily evaluated for all inputs u . For example the standard Gaussian where $u \sim \mathcal{N}(0, I)$ is a common choice. This means that f takes an input u , which follows the base density distribution $u \sim \pi_u(u)$ and transforms it to data space via the invertible function $x = f(u)$. As f is invertible it holds that $u = f^{-1}(x)$, and the density $p(x)$ can be calculated as

$$p(x) = \pi_u(f^{-1}(x)) \left| \det \left(\frac{\partial f^{-1}}{\partial x} \right) \right|. \quad (2.2)$$

By construction the computation of this expression is tractable. The first part is easy to compute because f is by design invertible and the base density can be evaluated for any input. The absolute determinant of the Jacobian is tractable, because the Jacobian matrix is triangular. This is due to the autoregressive prop-

erty where the d -th variable x_d only depends on the prior $d - 1$ variables $x_{1:d-1}$. Consequently, the Jacobian of f^{-1} is zero for all variables that this variable does not depend on. In each layer the mean and the log standard deviation of the d -th conditional given only the previous variables are computed via functions which are implemented as MADE units. Stacking multiple such units on top of each other is possible because if f_1 and f_2 are differentiable and invertible functions their composition $f_1 \circ f_2$ is also invertible and differentiable. This means the necessary properties to keep expression 2.2 tractable are maintained, even when stacking multiple layers. The use of masking makes it possible to compute the density $p(x)$ of data x in a single forward pass.

2.2 Introduction to Computational State Space Models

The goal of human activity recognition is to recognize activities in a sequence of potential actions that a human could perform. As such the HAR context is always a dynamic one, i.e. entities within this system are subject to changing states. Starting from some initial state and following a transition model the system state changes as time progresses. Due to combinatorial explosion the space of potential states is very large. Thus, a system which allows tractable inference in dynamic systems is required. In a feasibility study, Computational State Space models (CSSMs) have been shown to fulfill this property, even in complex, real-world applications [6]. The Authors found that the domain model used in their evaluation of the carrots data set had circa 10^8 states. Yet, the system was able to outperform Hidden Markov Models (HMMs) if the correct inference procedures were chosen. This section introduces the necessary background and notation to understand such models.

In CSSMs the transition model of the dynamic system is described by a computable function, which differentiates it from systems that require the explicit enumeration of states or paths, like HMMs. The behavior of the dynamic system is characterized as a labeled transition system (LTS). It consists of a set of states \mathcal{S} , a set of actions \mathcal{A} , and labeled transition relations $\rightarrow \subseteq \mathcal{S} \times \mathcal{A} \times \mathcal{S}$ that describe transitions between states. Thus, an LTS is a triple of the form $(\mathcal{S}, \mathcal{A}, \rightarrow)$. It is possible that \mathcal{S} and \rightarrow become infinite while \mathcal{A} is finite. As such it is necessary to find a suitable computational description to avoid explicit enumeration of all (potentially infinitely many) states. This description is then called the *computational action language* and is used to define the transition model. It is especially applicable when the modeled process (like cooking carrot soup) can be interpreted as performing sequential computations to arrive at some goal state starting from an initial state. That is, when the observed behavior is somehow goal directed.

In general the model should allow tractable inference in the sense that the hidden state of the LTS can be inferred at time t given a sequence of past observations $Y_{1:t}$ from a space \mathcal{Y} and a sequence of states $X_{1:t}$ from space \mathcal{X} . If the joint distribution $p(x_{1:t}, y_{1:t})$ factorizes over time into a *transition model* $p(x_t|x_{t-1})$ and an *observation model* $p(y_t|x_t)$ such that

$$p(x_{1:t}, y_{1:t}) = p(y_1|x_1)p(x_1) \prod_{i=2}^t (p(y_i|x_i)p(x_i|x_{i-1})), \quad (2.3)$$

the joint distribution can be described by a state space model. The underlying idea of this thesis is to obtain the observation model via a neural density estimator like MAF as outlined in the section 2.1.

Chapter 3

Construction of the Hybrid Model

Chapter 4

Experimental Evaluation

4.1 Settings

4.2 Experiments

4.3 Results

Chapter 5

Conclusion

5.1 Summary

5.2 Future Work

Bibliography

- [1] Luca Arrotta, Gabriele Civitarese, and Claudio Bettini. Knowledge Infusion for Context-Aware Sensor-Based Human Activity Recognition. In *2022 IEEE International Conference on Smart Computing (SMARTCOMP)*, pages 1–8, June 2022. ISSN: 2693-8340.
- [2] Kaixuan Chen, Dalin Zhang, Lina Yao, Bin Guo, Zhiwen Yu, and Yunhao Liu. Deep Learning for Sensor-based Human Activity Recognition: Overview, Challenges, and Opportunities. *ACM Computing Surveys*, 54(4):1–40, May 2022.
- [3] Liming Chen, Jesse Hoey, Chris D. Nugent, Diane J. Cook, and Zhiwen Yu. Sensor-Based Activity Recognition. *IEEE Transactions on Systems, Man, and Cybernetics, Part C (Applications and Reviews)*, 42(6):790–808, November 2012. Conference Name: IEEE Transactions on Systems, Man, and Cybernetics, Part C (Applications and Reviews).
- [4] Mathieu Germain, Karol Gregor, Iain Murray, and Hugo Larochelle. MADE: Masked Autoencoder for Distribution Estimation. page 9.
- [5] Ivan Kobyzev, Simon J.D. Prince, and Marcus A. Brubaker. Normalizing Flows: An Introduction and Review of Current Methods. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 43(11):3964–3979, November 2021. Conference Name: IEEE Transactions on Pattern Analysis and Machine Intelligence.
- [6] Frank Krüger, Martin Nyolt, Kristina Yordanova, Albert Hein, and Thomas Kirste. Computational State Space Models for Activity and Intention Recognition. A Feasibility Study. *PLOS ONE*, 9(11):e109381, November 2014. Publisher: Public Library of Science.
- [7] Stefan Lüdtke, Fernando Moya Rueda, Waqas Ahmed, Gernot A. Fink, and Thomas Kirste. Human Activity Recognition using Attribute-Based Neural

- Networks and Context Information, October 2021. arXiv:2111.04564 [cs, eess].
- [8] Robin Manhaeve, Sebastijan Dumancic, Angelika Kimmig, Thomas De-meester, and Luc De Raedt. DeepProbLog: Neural Probabilistic Logic Programming. In *Advances in Neural Information Processing Systems*, volume 31. Curran Associates, Inc., 2018.
 - [9] Fernando Moya Rueda, René Grzeszick, Gernot A. Fink, Sascha Feldhorst, and Michael Ten Hompel. Convolutional Neural Networks for Human Activity Recognition Using Body-Worn Sensors. *Informatics*, 5(2):26, June 2018. Number: 2 Publisher: Multidisciplinary Digital Publishing Institute.
 - [10] George Papamakarios, Theo Pavlakou, and Iain Murray. Masked Autoregressive Flow for Density Estimation. In *Advances in Neural Information Processing Systems*, volume 30. Curran Associates, Inc., 2017.
 - [11] Luc De Raedt. ProbLog: A Probabilistic Prolog and Its Application in Link Discovery. page 6.
 - [12] Miquel Ramirez and Hector Geffner. Goal Recognition over POMDPs: Inferring the Intention of a POMDP Agent. In *Twenty-second international joint conference on artificial intelligence*, page 6, 2011.
 - [13] Fernando Moya Rueda, Stefan Lüdtke, Max Schröder, Kristina Yordanova, Thomas Kirste, and Gernot A. Fink. Combining Symbolic Reasoning and Deep Learning for Human Activity Recognition. In *2019 IEEE International Conference on Pervasive Computing and Communications Workshops (Per-Com Workshops)*, pages 22–27, March 2019.
 - [14] Yoli Shavit and Itzik Klein. Boosting Inertial-Based Human Activity Recognition With Transformers. *IEEE Access*, 9:53540–53547, 2021. Conference Name: IEEE Access.

Appendix A

Program Code / Resources

The source code, a documentation, some usage examples, and additional test results are available at ...

They as well as a PDF version of this thesis is also contained on the CD-ROM attached to this thesis.

Appendix B

Further Experimental Results

In the following further experimental results are ...

Ehrenwörtliche Erklärung

Ich versichere, dass ich die beiliegende Master-/Bachelorarbeit ohne Hilfe Dritter und ohne Benutzung anderer als der angegebenen Quellen und Hilfsmittel angefertigt und die den benutzten Quellen wörtlich oder inhaltlich entnommenen Stellen als solche kenntlich gemacht habe. Diese Arbeit hat in gleicher oder ähnlicher Form noch keiner Prüfungsbehörde vorgelegen. Ich bin mir bewusst, dass eine falsche Erklärung rechtliche Folgen haben wird.

Mannheim, den 31.09.2022

Unterschrift